

Our framework vs published-method baselines (Self-Refine, Reflexion) — Δ utility vs always_verify
Bayesian wins 19/20 cells. Reflexion wins on LCB-easy gpt5_mini (regime crossover: high prior + P_fix=0.53).

