

our framework Self-Refine [Madaan 2023] Reflexion [Shinn 2023]
Our framework vs published-method baselines (Δ utility vs `always_verify`)
Bayesian wins 19/20 cells. One crossover: LCB-easy/gpt5-mini (high prior + high P_{fix})

