

Bayesian POMDP Controllers for LLM Critic Stack Orchestration

Research progress & experimental results

MBZUAI

May 8, 2026

Outline

Click any item to jump to that section. Colors: **blue**=core conceptual blocks, **green**=positive findings, **orange**=notable / has caveat, **red**=limitations, **teal**=cross-reference.

1. **The problem we're solving**
2. Existing methods (**baselines**)
3. How baselines work on real data — **examples**
4. **Statistics we collect** (priors, likelihoods, kernels)
5. **Our proposed method**: Bayesian POMDP controller
6. **Benchmarks** (LCB, SWE-bench, MBPP+, HumanEval+)
7. Training-data collection pipeline
8. Iterative refinement (Self-Refine, Reflexion) — **real implementations**
9. Evaluation methodology (bootstrap CI, FDR, MDE-power, LOO-CV)
10. **Pre-registered prediction** & confirmation
11. **Three-regime characterization**
12. Results across 80+ cells
13. **Honest negative findings**
14. Open-weight integration (Qwen2.5-Coder-32B)
15. **Limitations** & next steps

What we're doing — one-sentence version

Research question

Given an LLM that proposes a code patch and a stack of cheap-to-expensive critics (syntax check, lint, public tests, LLM review), **when should we verify a patch with a costly oracle vs. regenerate?**

Why this matters:

- ▶ **Verification is the dominant cost** — running full test suites, calling frontier LLMs as judges, or paying for human review.
- ▶ **Most pipelines use rigid heuristics:** “always verify,” “best-of- N ,” “run all critics.” **They ignore the posterior probability** that the patch is correct given cheap signals.
- ▶ **A principled answer is a partially-observable MDP:** state = $P(Y=1 \mid \text{evidence})$, actions $\in \{\text{run critic } i, \text{verify, stop, regen}\}$, solve via dynamic programming.

Concrete setup

Per LLM patch (instance i , attempt k):

1. Critics produce noisy binary signals $L_0, L_1, L_2, L_3 \in \{0, 1\}$ (syntax, lint, public-tests, LLM-review).
2. Oracle verifier produces ground truth $Y \in \{0, 1\}$ (full test suite).
3. Each critic has a known cost c_0, c_1, c_2, c_3 , oracle cost c_v .
4. A reward r is paid for $Y=1$.

Optimal policy maximises $\mathbb{E}[r \cdot Y - \text{costs}]$ over policies $\pi : \text{history} \rightarrow \text{action}$.

Two regimes our framework distinguishes

- ▶ **Inference-driven:** cheap critics shift the posterior enough to skip the oracle on confident negatives. Bayesian controller wins.
- ▶ **Oracle-cheap:** verifying is so cheap that always-verify is indistinguishable. Bayesian gain ≈ 0 .

Baselines we compare against

Policy	What it does
always_verify	Run oracle on every first patch.
best_of_N	Generate N patches, verify last (or vote).
threshold_L0	Run cheap syntax check. Verify iff PASS. Else regen.
threshold_L2	Run public tests. Verify iff PASS. Else regen. (LCB-only)
threshold_L3	Run LLM-judge review. Verify iff PASS. Else regen.
fixed_pipeline	Run <i>all</i> cheap critics, verify iff all PASS.
Iterative refinement methods	
Self-Refine (Madaan 2023)	Generate $\rightarrow$ self-critique $\rightarrow$ refine. Stops on "CRITIQUE_OK."
Reflexion (Shinn 2023)	Generate $\rightarrow$ external test $\rightarrow$ reflect $\rightarrow$ refine. Stops on test pass.

First 6: **stateless per-patch policies**. **Self-Refine** and **Reflexion** are stateful trajectories we treat as competing methods.

Our methods

Method	What it does
Bayesian belief-state controllers (operate on $b_t = P(Y_t=1 \mid \text{history})$)	
<code>bayesian_greedy</code>	1-step lookahead: at each state, pick argmax of immediate Q-value. Doesn't use the kernel; treats one critic + verify as the full plan.
<code>bayesian_DP</code>	Backward induction on discretised belief grid ($51 \text{ pts} \times \text{horizon } H=3$). Uses the <i>transition kernel</i> $P(Y_{k+1} \mid Y_k)$; can chain critics multi-step.

Shared by both:

- ▶ Same belief update

$$b' = \frac{P(z|Y=1)b}{P(z|Y=1)b + P(z|Y=0)(1-b)}$$
- ▶ Same action set

$$\{\text{verify, generate, give_up}, L_0, L_2, L_3\}$$
- ▶ Same cost vector $(c_{\text{gen}}, c_{\text{ver}}, c_{L_0}, c_{L_2}, c_{L_3})$

Differ on:

- ▶ Planning horizon: 1 vs $H=3$
- ▶ Uses kernel for planning? No vs Yes
- ▶ Can chain critics? No vs Yes
- ▶ Robust to kernel misspecification? More so vs less

Code: `GreedyController` in `lcb_compare.py`:52 `BayesianController (DP)` in

`run_baseline_vs_controller.py`:68.

bayesian_greedy — 1-step lookahead

Mechanism: at every state (b, k) , evaluate each action's *immediate* Q -value assuming the very next move is verify. **Does not use the transition kernel** — by design, Q_{generate} substitutes the prior for the post-generate belief.

What data this method needs

Calibration corpus only: prior $P(Y=1)$ + critic likelihoods $P(z | Y)$.

No iter / no kernel needed for setup. Iter trajectories are only used later, at evaluation, to simulate the policy and compute utility.

+ a cost vector $(c_{\text{gen}}, c_{\text{ver}}, c_{L_0}, c_{L_2}, c_{L_3})$.

Q-values (myopic, 1-step horizon)

$$Q_{\text{verify}} = R \cdot b - c_{\text{ver}}$$

$$Q_{\text{generate}} = -c_{\text{gen}} + R \cdot \text{prior} - c_{\text{ver}} \quad (\text{uses prior, NOT kernel})$$

$$Q_{L_k} = -c_k + \mathbb{E}_Z [\max(R \cdot b' - c_{\text{ver}}, 0)]$$

Worked example — LCB-medium / haiku-4.5 (prior 0.34, L_2 gap 0.92, $R=100$)

Initial $b_0 = 0.34$: $Q_{\text{verify}}=33.0$, $Q_{L_2}=37.3$, $Q_{\text{give_up}}=0$.
 $\Rightarrow \text{argmax} = L_2$. Observe $z_2=\text{PASS} \Rightarrow b_1 \approx 0.97 \Rightarrow$ next step picks **verify**.

bayesian_DP — backward-induction controller

Mechanism: discretise belief into 51 grid points $\times$ horizon $H=3$; solve Bellman backwards; act greedily w.r.t. the resulting V-table. **Uses the measured kernel** $P(Y_{k+1} | Y_k)$ for generate; recurses on critic outcomes (can plan $L0 \rightarrow L2 \rightarrow \text{verify}$ multi-step).

What data this method needs

Calibration corpus $\rightarrow$ prior + likelihoods (same as greedy). **Iter corpus** (multi-step Self-Refine / Reflexion) $\rightarrow$ measured kernel via `compute_transition_kernel.py`, giving $P(\text{fix} | \text{broken})$, $P(\text{break} | \text{correct})$ from $Y_k \rightarrow Y_{k+1}$ pairs.

Bellman backward induction (build once per cell)

$$V_H(b) = \max(Rb - c_{\text{ver}}, 0); \quad Q_k(a; b) = \begin{cases} Rb - c_{\text{ver}} & a=\text{verify} \\ -c_{\text{gen}} + V_{k+1}(\text{kernel}(b)) & a=\text{generate} \\ -c_a + \mathbb{E}_z[V_k^{\text{post}}(b'_z)] & a=L_0, L_2, L_3 \\ 0 & a=\text{give_up} \end{cases}$$

Worked example — LCB-easy / gpt-5-mini ($P_{\text{fix}}=0.53$, $P_{\text{break}}=0.04$)

$b_0=0.22$. **Greedy** substitutes prior $\Rightarrow Q_{\text{generate}} = 20$. **DP** applies kernel:
 $\text{kernel}(b_0) = 0.22 \cdot 0.96 + 0.78 \cdot 0.53 = \mathbf{0.625}$, $Q_{\text{generate}} = -1 + V_1(0.625) = \mathbf{59.4} \Rightarrow \text{generate} \rightarrow L_2$ at
 $b_1=0.625 \rightarrow \text{verify}$ at $b_2 \approx 0.99$.

Self-Refine vs Reflexion — algorithmically

Self-Refine (no external feedback)

```
patch_0 = generate(prompt)
for k = 1..N:
    critique = model.self_critique(patch_{k-1})
    if "CRITIQUE_OK" in critique:
        return patch_{k-1}
    patch_k = model.refine(patch_{k-1}, critique)
return patch_N
```

Reflexion (external test signal)

```
memory = []
patch_0 = generate(prompt)
for k = 1..N:
    pass, trace = external_test(patch_{k-1})
    if pass: return patch_{k-1}
    reflection = model.reflect(patch, trace, memory)
    memory.append(reflection)
    patch_k = model.refine(patch, memory)
return patch_N
```

Example: gpt5_mini on astropy__astropy-14182 — clean

Problem (excerpt): *add support for the `header_rows` argument in the `RST` writer to allow customising header rows.*

LLM patch (clean SEARCH/REPLACE format):

```
--- a/astropy/io/ascii/rst.py
+++ b/astropy/io/ascii/rst.py
@@ -57,10 +57,15 @@
     data_class = SimpleRSTData
     header_class = SimpleRSTHeader

-    def __init__(self):
+    def __init__(self, header_rows=None):
         super().__init__(delimi...
```

Critic verdicts on this patch:

- ▶ L0_syntax (ast.parse): **PASS**
- ▶ L1_lint (ruff): **PASS**
- ▶ L3_llm_review (Haiku-4.5): **PASS**
- ▶ Y (full test suite): **1 (resolved)**

All critics correctly flagged this as a likely-correct patch.

Example: haiku45 on same task — hallucination

Same problem as before.

LLM patch:

```
--- a/astropy/io/ascii/rst.py
+++ b/astropy/io/ascii/rst.py
@@ -57,10 +57,15 @@
     data_class = SimpleRSTData
     header_class = SimpleRSTHeader

-    def __init__(self, header_rows=None):
+    def __init__(self, header_rows=None):
+
+        # <-- HALLUCINATED!
+        super().__init__(delimi...
```

Issue: the SEARCH block claims the original code already had `header_rows=None`. Real source has `def __init__(self):`. Patch fails to apply.

Critic verdicts:

- ▶ L0_syntax: **PASS** (the diff is syntactically valid Python)
- ▶ L1_lint: **PASS**
- ▶ L3_llm_review: **PASS** (judge can't see real source)
- ▶ Y: **0 (unresolved — patch didn't apply!)**

All cheap critics fooled. This is why we need an oracle in the loop.

Example: sonnet45 on django__django-11620 — format quirk

LLM output (single-bracket header instead of triple —/+++ format):

```
[django/db/models/sql/query.py]
<<<<<<< SEARCH
def deferred_to_data(self, target, callback):
    ...
=====
def deferred_to_data(self, target, callback):
    callback(target, self)
    for ...
>>>>>>> REPLACE
```

Issue: format works via our XML-style parser fallback but is non-canonical. This is why we landed an **XML-style parser + fuzzy SEARCH/REPLACE matcher** (commit c10eb8a): handles

- ▶ placeholder paths ([file.py]) on the line before the diff
- ▶ indent-tolerant SEARCH/REPLACE (whitespace doesn't have to be byte-identical)
- ▶ triple-bracket and double-bracket variants

Validation: v3 parser pass rate from ~60% (canonical only) to ~93% (canonical + tolerant) on the spot-check $n = 50$ corpus.

Three quantities are needed for the controller

1. Prior $P(Y=1)$

Per-(generator, benchmark, difficulty) cell, fraction of step-0 patches that *would* resolve the bug under the oracle.

2. Critic likelihoods $P(\text{critic} = z \mid Y)$

For each critic $\ell \in \{L_0, L_1, L_2, L_3\}$, the conditional pass rate under both ground-truth values. Estimated with Beta(1,1) Laplace smoothing on the calibration corpus.

3. Transition kernel $P(Y_{k+1} \mid Y_k)$

$P(\text{fix} \mid \text{broken})$, $P(\text{break} \mid \text{correct})$ across iterative refinement steps. Measured directly from the iter trajectory data — not assumed.

Two modes of data collection: **calibration** vs **iter**

We collect data in **two distinct modes**, each producing a different kind of evidence about the (generator, benchmark) cell. Both are needed.

Mode 1: **Calibration mode** (single-shot, independent attempts)

- ▶ Sample n instances; for each, generate $k = 3$ *independent fresh* patches.
- ▶ Per patch: run all critics + harness oracle to get Y .
- ▶ Captures the *generation distribution*: “how good is this LLM at this benchmark, fresh attempt, no feedback?”

Used to estimate quantities that require an unconditional sample: **prior** $P(Y=1)$, **critic likelihoods** $P(z \mid Y)$.

Mode 2: **Iter mode** (sequential trajectories with feedback)

- ▶ Same n instances; for each, run a Self-Refine or Reflexion trajectory of up to 5 steps.
- ▶ Step $k + 1$ is *conditioned on* step k 's output via the controller's feedback signal.

Two corpora, each for what it cleanly identifies

Corpus	Size per cell	Used for
Calibration (single-shot)	30 instances $\times$ 3 fresh patches = 90 step-0 patches	Prior $P(Y=1)$ Critic likelihoods $P(z Y)$
Iterative trajectories (Self-Refine, Reflexion)	30 instances $\times$ up to 5 steps $\approx$ 150 step-records	Transition kernel $P(Y_{k+1} Y_k)$

Each parameter is estimated from the source where it's cleanly identifiable.

At decision time, the controller uses all three simultaneously:

- ▶ **Prior** seeds belief b_0 at step 0
- ▶ **Likelihoods** update b as critic outcomes are observed
- ▶ **Kernel** propagates b across refinement steps

Standard partially-observable MDP setup. Same pattern as forecasting: unconditional climatology + observation likelihoods + transition dynamics.

Why single-shot data for prior + likelihoods

The selection-bias trap:

- ▶ Step 1+ patches in iter trajectories are *only* drawn from instances where step 0 failed (Self-Refine/Reflexion continued because the proxy stop condition wasn't met).
- ▶ Including step 1+ patches when estimating the prior would **systematically bias it downward** — those are conditioned on prior failure, not a fresh attempt.
- ▶ Same argument for critic likelihoods: $P(z \mid Y)$ for a multi-step patch is not the same distribution as for a fresh step-0 patch.

Single-shot calibration gives a clean unconditional measurement of the *generation* distribution.

Forecasting analogy

Climatology measures unconditional $P(\text{rain})$ from independent single-day observations. You don't measure it from conditional rates at later forecast horizons — those depend on what already happened. Same logic here.

Recipe 1a: prior $P(Y=1)$ — data collection

Goal: per (generator, benchmark, difficulty) cell, the fraction of step-0 patches that resolve the bug under the oracle.

Steps 1–3 (collecting raw Y values):

1. **Sample $n = 30$ instances** from the benchmark with seed=42. Same instance set across all generators for reproducibility.
2. **For each instance, generate $k = 3$ candidate patches** via the LLM (temperature 0.7, distinct system messages):
 - ▶ OpenRouter API for closed-API generators
 - ▶ local vLLM endpoint for Qwen2.5-Coder-32B

Output $\rightarrow$ predictions_p{0,1,2}.jsonl (~ 90 patches/cell).

3. **Run harness eval** on every patch:
 - ▶ LCB: hidden-test runner returns pass/fail
 - ▶ SWE: podman container, applies patch, runs FAIL_TO_PASS + PASS_TO_PASS test sets
 - ▶ MBPP+/HumanEval+: assert-based unit tests

Record $Y_i \in \{0,1\}$ per patch into critic_results.jsonl.

Recipe 1b: prior $P(Y=1)$ — estimator & CI

Steps 4–6 (computing the prior):

4. **Compute the point estimator** (sample mean):

$$\hat{P}(Y=1) = \frac{1}{N} \sum_{i=1}^N Y_i, \quad N = \#\{\text{patches with non-empty diff}\}$$

5. **Wilson 95% confidence interval** on the proportion (better than normal-approx for small N near 0/1 boundaries):

$$\text{CI}_{95\%} = \frac{\hat{p} + \frac{z^2}{2N} \pm z \sqrt{\frac{\hat{p}(1-\hat{p})}{N} + \frac{z^2}{4N^2}}}{1 + \frac{z^2}{N}}, \quad z = 1.96$$

6. **Persist** to `likelihood_tables.json`: `{"prior_Y1": $\hat{p}$,
"prior_Y1_ci": [lo, hi], "n": N}`

Concrete example

LCB-hard / gpt5_mini: 18 of 87 non-empty patches resolved
 $\Rightarrow \hat{P}(Y=1) = \mathbf{0.21}$, Wilson CI $\mathbf{[0.13, 0.31]}$

Recipe 2a: critic likelihoods $P(z \mid Y)$ — verdicts

Goal: $P(\text{critic } \ell \text{ passes} \mid Y=1)$ vs. $P(\text{passes} \mid Y=0)$ for $\ell \in \{L_0, L_1, L_2, L_3\}$.

Notation: $z_\ell \in \{0, 1\}$ is the binary verdict of critic ℓ on a patch:

- ▶ $z_\ell = 1 \Leftrightarrow$ critic outputs **PASS** (claims patch is correct).
- ▶ $z_\ell = 0 \Leftrightarrow$ critic outputs **FAIL** (claims patch is incorrect).

$Y \in \{0, 1\}$ is the oracle ground truth from the harness.

Steps 1–3 (per-patch verdicts):

1. **Reuse the same 90 patches** from Recipe 1, with Y already known.

2. **Compute each critic's verdict per patch:**

- ▶ L_0 **syntax**: `ast.parse(modified_file)` for each `.py` touched by the diff $\rightarrow$ True iff all parse cleanly.
- ▶ L_1 **lint**: `ruff check` on the modified files $\rightarrow$ True iff zero diagnostics.
- ▶ L_2 **public tests** (LCB-only): run the public examples shipped with the problem $\rightarrow$ True iff all pass.
- ▶ L_3 **LLM review**: prompt Haiku-4.5 with diff + problem statement, parse one-token PASS/FAIL response.

3. **Build 2×2 confusion matrix** per critic against ground truth Y :

$$\text{TP}_\ell = \#\{z_\ell=1 \wedge Y=1\}, \quad \text{FN}_\ell = \#\{z_\ell=0 \wedge Y=1\}$$

$$\text{FP}_\ell = \#\{z_\ell=1 \wedge Y=0\}, \quad \text{TN}_\ell = \#\{z_\ell=0 \wedge Y=0\}$$

Recipe 2b: critic likelihoods $P(z | Y)$ — smoothing & tests

Steps 4–7 (estimating + validating):

4. **Beta(1,1) Laplace smoothing** (avoids 0/1 estimates with small N):

$$\hat{P}(z_\ell=1 | Y=1) = \frac{TP_\ell + 1}{TP_\ell + FN_\ell + 2}, \quad \hat{P}(z_\ell=1 | Y=0) = \frac{FP_\ell + 1}{FP_\ell + TN_\ell + 2}$$

5. **Compute the gap** (per-critic informativeness):

$$\Delta_\ell = \hat{P}(z_\ell=1 | Y=1) - \hat{P}(z_\ell=1 | Y=0)$$

Higher gap $\Rightarrow$ more informative critic.

6. **χ^2 critic-independence test:** verify $L_0 \perp L_1 \perp L_3 | Y$ assumption (we assume *product* likelihood for the Bayesian update):

$$\chi^2 = \sum_{\text{cells}} \frac{(O_{ij} - E_{ij})^2}{E_{ij}}$$

$p > 0.05$ on all cells $\Rightarrow$ independence assumption safe.

7. **Persist** to `likelihood_tables.json`: per-critic $\{TP, FN, FP, TN, \hat{P}(\cdot | Y=1), \hat{P}(\cdot | Y=0), \text{gap}\} + \chi^2$ test summary.

Example LCB-hard gpt5_mini: L_2 gap = +0.49 (highly informative), L_3 gap = +0.07 (mild).

Recipe 3a: transition kernel $P(Y_{k+1} \mid Y_k)$ — collect

Goal: measured Markov dynamics across iterative refinement steps.

Steps 1–4 (collecting trajectories & classifying):

1. **Run iterative refinement** (Self-Refine or Reflexion) on the same $n = 30$ instances used for the prior. For each instance, record up to 5 trajectory steps with method-specific stop conditions.
2. **Persist** per-step row to `iter_records.jsonl`:

```
{
  "instance_id": "django_django-12284",
  "step": 1,
  "diff": "...",
  "Y": 1,
  "L0": true,
  "L1": true,
  "L3": true,
  "stop_reason": "...",
  "step_cost_usd": 0.012
}
```

3. **Group by** `instance_id` and sort within each group by step.
4. **For each consecutive pair** (Y_k, Y_{k+1}) within an instance, classify into 4 categories:
 - ▶ $Y_k=0, Y_{k+1}=1$: **fix** (broken $\rightarrow$ correct)
 - ▶ $Y_k=0, Y_{k+1}=0$: **persist_broken**
 - ▶ $Y_k=1, Y_{k+1}=0$: **break** (correct $\rightarrow$ broken)
 - ▶ $Y_k=1, Y_{k+1}=1$: **persist_correct**

Recipe 3b: transition kernel $P(Y_{k+1} | Y_k)$ — estimate

Steps 5–8 (estimating, bootstrap, validation):

5. Estimators (aggregated across all consecutive pairs in the cell):

$$\hat{P}(\text{fix} | \text{broken}) = \frac{\# \text{fix}}{\# \text{fix} + \# \text{persist_broken}}$$

$$\hat{P}(\text{break} | \text{correct}) = \frac{\# \text{break}}{\# \text{break} + \# \text{persist_correct}}$$

6. Cluster bootstrap at the *instance level* ($B=1000$):

- ▶ Resample whole-instance trajectories with replacement
- ▶ Recompute the kernel per resample
- ▶ Empirical 2.5/97.5 percentiles $\rightarrow$ CI

Why cluster, not row: within-instance steps auto-correlate (step k 's diff depends on step $k-1$). Per-row bootstrap underestimates SE by $\sim 40\%$.

7. Compare against IID baseline: $\hat{P}_{\text{IID}}(\text{fix}) = \hat{P}(Y=1)$ (the prior). **If measured $\neq$ IID, iid assumption is wrong** $\Rightarrow$ DP solver may strictly dominate Greedy. (§F1 ablation confirms this on LCB-hard.)

8. Persist to `transition_kernel.json`: {measured kernel, IID baseline kernel, cluster-bootstrap CI, per-step counts}.

Example gpt5_mini SWE-Verified: $\hat{P}(\text{fix}|\text{broken}) = 0.13$, $\hat{P}(\text{break}|\text{correct}) = 0.06$. Markov chain $\neq$ iid

Likelihood table — concrete example

LCB-hard, gpt5_mini, $n = 87$ patches across 29 instances:

Critic	$P(z \mid Y=1)$	$P(z \mid Y=0)$	gap Δ
L0_syntax	0.95	0.99	-0.04
L1_lint	0.95	0.99	-0.04
L2_public_tests	0.90	0.41	+0.49
L3_llm_review	0.30	0.23	+0.07

- ▶ **L0/L1 are uninformative for Y** — almost all patches parse cleanly.
- ▶ **L2 (public tests) is highly informative** — gap 0.49 is the workhorse signal for LCB.
- ▶ **L3 (LLM judge) is mildly informative** — gap 0.07.
- ▶ **Bayesian update:** $\frac{P(Y=1)}{P(Y=0)} \cdot \prod_{\ell} \frac{P(z_{\ell}|Y=1)}{P(z_{\ell}|Y=0)}$.

What does the transition kernel depend on?

One kernel per (generator, benchmark, difficulty, refinement method) cell:

$$P(Y_{k+1} \mid Y_k; \text{gen, bench, diff, method})$$

Dependencies:

- ▶ **Generator** — different LLMs refine differently (gpt5_mini's $P(\text{fix}|\text{broken}) \neq \text{qwen3's}$)
- ▶ **Benchmark + difficulty** — LCB-hard problems are intrinsically harder to fix than LCB-easy
- ▶ **Refinement method** — Self-Refine (model self-critique) and Reflexion (external test) yield *different* kernels because the feedback signal differs

What it does NOT depend on:

- ▶ **Not per-critic** — the kernel describes the transition of the *oracle outcome* Y , not how critics correlate with Y
- ▶ **Not per-cost-vector** — the dynamics are intrinsic; cost only matters at decision time

For our matrix: 5 generators $\times$ {LCB $\times$ 3, SWE-Lite, SWE-Verified} $\times$ 2 methods = **50 distinct kernels** estimated.

Why kernel and critics are estimated independently

Conceptual factorization (the Bayesian controller's joint over a step):

$$\underbrace{P(Y_{k+1}, z_{k+1}^{(\ell)} \mid Y_k)}_{\text{joint}} = \underbrace{P(Y_{k+1} \mid Y_k)}_{\text{kernel}} \cdot \underbrace{P(z_{k+1}^{(\ell)} \mid Y_{k+1})}_{\text{critic likelihood}}$$

Two cleanly separated quantities:

- ▶ **Kernel** answers: “given the model’s last patch was broken, what’s the chance the refined patch is correct?”
 - ▶ *Source of randomness:* model’s stochasticity, problem difficulty
 - ▶ *Estimated from:* iter trajectories
- ▶ **Critic likelihood** answers: “given a patch is correct (or broken), what’s the probability the critic says PASS?”
 - ▶ *Source of randomness:* critic’s own noise (LLM-judge stochasticity, lint sensitivity, etc.)
 - ▶ *Estimated from:* single-shot calibration

Critical assumption: critics are *conditionally independent of the trajectory history given Y_{k+1}* . We verify this with the χ^2 independence test (Recipe 2b).

This is why we need only 50 kernels (not 50×4 critics = 200): the critic conditional structure is already captured by the likelihoods.

Concrete worked kernel (1/2): counts & estimators

Cell: verified/gpt5_mini/selfrefine, $n = 30$ instances, trajectories of up to 5 steps each.

Step 1: count consecutive transitions across all instances

Transition	Y_k	Y_{k+1}	Count
fix	0	1	18
persist_broken	0	0	122
break	1	0	9
persist_correct	1	1	130

Step 2: maximum-likelihood estimators

$$\hat{P}(\text{fix} \mid \text{broken}) = \frac{18}{18 + 122} = \mathbf{0.13}, \quad \hat{P}(\text{break} \mid \text{correct}) = \frac{9}{9 + 130} = \mathbf{0.06}$$

Next slide: cluster-bootstrap CI on these estimates and the IID-baseline comparison that justifies using a measured kernel rather than a synthesised one.

Concrete worked kernel (2/2): CI & IID-baseline comparison

Same cell: verified/gpt5_mini/selfrefine. Point estimates from the previous slide: $\hat{P}(\text{fix} \mid \text{broken}) = \mathbf{0.13}$, $\hat{P}(\text{break} \mid \text{correct}) = \mathbf{0.06}$.

Step 3: cluster bootstrap CI ($B = 1000$, instance-level)

- ▶ Resample 30 instances with replacement, recompute kernel per resample.
- ▶ Clustering at the instance level handles within-trajectory auto-correlation (step $k + 1$ conditions on step k).
- ▶ 95% CI: $P(\text{fix} \mid \text{broken}) \in \mathbf{[0.07, 0.20]}$, $P(\text{break} \mid \text{correct}) \in \mathbf{[0.02, 0.11]}$.

Step 4: compare to IID baseline (assuming patches are iid Bernoulli(prior))

- ▶ IID would predict $P(\text{fix} \mid \text{broken}) = \text{prior} = 0.63$.
- ▶ **Measured 0.13 $\ll$ IID 0.63 $\Rightarrow$ IID assumption strongly violated.**
- ▶ Implication: DP solver may strictly dominate Greedy here (§F1 ablation).

Transition kernel — measured vs IID synthesis

Naive assumption (IID): per-step patches are independent draws from Bernoulli(prior).

Reality (measured kernels on iter records):


- ▶ $P(\text{fix} \mid \text{broken}) = 0.13$ (gpt5_mini SWE-Verified)
- ▶ $P(\text{break} \mid \text{correct}) = 0.06$
- ▶ **Markov chain $\neq$ iid Bernoulli.**

Verification across the cube:

- ▶ **IID assumption is empirically wrong on every benchmark** (commit ef21a0a).
- ▶ **DP $\succ$ Greedy in some regimes** when measured kernel is used (§F1 ablation).
- ▶ Below LCB-hard, Greedy $\approx$ DP — kernel non-iid'ness doesn't matter much because pass rates saturate.

Transition kernel: visual — iter feedback at full LCB sample

Iterative refinement vs IID-baseline retries (LCB, full sample $n=102/207/135$)
 $\text{iter} \leq \text{IID baseline}$ for nearly all closed-API cells; only qwen32b/LCB-hard shows positive Δ



Top: measured $P(\text{fix}|\text{broken})$ (solid bar) vs IID baseline (hatched, computed from $(\text{instance}, \text{patch}; d)$ pairs in calibration). **Bottom:** $\Delta = \text{iter} - \text{IID}$. **Honest finding at full LCB sample ($n=102/207/135$):** $\text{iter} \leq \text{IID baseline}$ for nearly all closed-API cells. Only qwen32b/LCB-hard shows positive Δ (+0.007). The earlier “iter helps capable gens on easy” story was a small-sample artefact.

Our proposed method — POMDP formulation

Belief state: $b_k = P(Y_k=1 \mid \text{history of critic outcomes})$, maintained via Bayes' rule.

Action space at each timestep:

- ▶ Run critic ℓ (cost c_ℓ , observe z_ℓ , update b)
- ▶ Verify (cost c_v , observe Y , terminate)
- ▶ Regenerate (cost c_g , draw a fresh patch, reset b to prior)

Two solvers:

1. **Greedy / myopic** — at each step pick action that maximises one-step expected reward.
2. **DP / backward-induction** — solve the Bellman equation explicitly, optimal in the iid regime, near-optimal under the measured Markov kernel.

F1 ablation: **Greedy and DP coincide** for most cells; **DP wins** on cells where the kernel is meaningfully non-stationary.

Cost model and reward

Default cost vector (units = “LLM-cents”):

$$c_{L_0} = 1, c_{L_1} = 1, c_{L_2} = 5, c_{L_3} = 5, c_v = 30, r = 100.$$

Sensitivity sweeps on c_v : cell $\rightarrow$ regime label changes when c_v crosses a threshold (commit 1e9b354).

- ▶ $c_v < \text{threshold}$: “always verify” is optimal — Bayesian gain $\rightarrow 0$.
- ▶ $c_v > \text{threshold}$: cheap-critic gating kicks in — Bayesian wins.

Why this matters for the paper: the win is conditional on the cost regime; we report results across $c_v \in \{15, 20, 25, 30, 40, 60\}$ to show robustness.

Utility — what we optimize

Per-instance utility under policy π on instance i :

$$U_{\pi}(i) = \underbrace{R \cdot Y_i \cdot \mathbb{I}[\text{verified}]}_{\text{reward only if verified \& } Y=1} - \underbrace{\sum_{a \in \text{actions}_{\pi}(i)} c_a}_{\text{sum of action costs}}$$

Per-cell mean utility:

$$\bar{U}_{\pi} = \frac{1}{N} \sum_{i=1}^N U_{\pi}(i)$$

Headline metric reported throughout the paper:

$$\Delta_{\pi} = \bar{U}_{\pi} - \bar{U}_{\text{always_verify}}$$

Worked example: gpt5_mini always_verify on LCB-hard ($n = 29$, prior 0.21)

$$\bar{U}_{\text{always_verify}} = \frac{1}{29} \sum_i (R \cdot Y_i - c_g - c_v) = 100 \cdot 0.21 - 5 - 30 = -14.$$

Bayesian-greedy: $\bar{U}_{\pi} = -14 + 12.55 = -1.45$. $\Delta = +12.55$. The Bayesian framework recovers ~ 13 utility units per instance by gating on cheap critics.

Step 0: where prior + likelihoods are used

Concrete instance: LCB-hard problem 2784, generator gpt5_mini.

Belief state initialisation:

$$b_0 = P(Y=1 \mid \emptyset) = \text{prior} = 0.21$$

Step A: run L_2 (public tests, cost 5)

- ▶ Observe $z_{L_2} = \text{PASS}$.
- ▶ Bayes' rule with $P(L_2\text{pass} \mid Y=1) = 0.90$, $P(L_2\text{pass} \mid Y=0) = 0.41$:

$$b_1 = \frac{0.21 \times 0.90}{0.21 \times 0.90 + 0.79 \times 0.41} = \frac{0.189}{0.513} = 0.37$$

Step B: run L_3 (LLM review, cost 5)

- ▶ Observe $z_{L_3} = \text{PASS}$.
- ▶ $P(L_3 \mid Y=1) = 0.30$, $P(L_3 \mid Y=0) = 0.23$:

$$b_2 = \frac{0.37 \times 0.30}{0.37 \times 0.30 + 0.63 \times 0.23} = \frac{0.111}{0.256} = 0.43$$

Decision: $b_2 = 0.43 < 0.5 \Rightarrow$ **don't verify**; regenerate.

$\Rightarrow$ Cost so far: $c_g + c_{L_2} + c_{L_3} = 5 + 5 + 5 = 15$. **Saves $c_v = 30$ vs `always_verify`.**

Step $k \rightarrow k+1$: where the transition kernel is used

Continuing from same instance: controller chose to regenerate at step 0. Now at step 1.

What changed? The fresh generation isn't drawn from the prior anymore — it's conditioned on *having just failed step 0*. We use the kernel to propagate the belief.

Kernel from iter trajectories on this cell:

- ▶ $P(\text{fix} \mid \text{broken}) = 0.13$ (gpt5_mini SWE-Verified, similar on LCB-hard)
- ▶ $P(\text{break} \mid \text{correct}) = 0.06$

Updated belief at step 1:

$$b_1^{(\text{step}1)} = \underbrace{(1 - b_0^{\text{post-step}0}) \cdot P(\text{fix} \mid \text{broken})}_{\text{transition from broken}} + \underbrace{b_0^{\text{post-step}0} \cdot (1 - P(\text{break} \mid \text{correct}))}_{\text{stay correct}}$$

Cost-benefit decision (refine vs. stop):

- ▶ Expected gain from refinement: $0.13 \times 100 = 13$ utility
- ▶ Refinement cost: $c_g = 5$
- ▶ **Net: $+8 \Rightarrow$ refine and reassess**

This decision is impossible without the kernel. A 1-step controller has no notion of “what's the chance refinement fixes it” — it'd just regenerate from the prior.

Side-by-side: the three statistics inside one trajectory

Single instance walkthrough: all three statistics appear at different decision points.

Step	What controller does	Stat used	Cost
0a	Initialise $b_0 = 0.21$	Prior (from calibration)	0
0b	L_2 pass $\rightarrow b = 0.37$	Likelihood $P(L_2 \mid Y)$ (calibration)	$c_{L_2} = 5$
0c	L_3 pass $\rightarrow b = 0.43$	Likelihood $P(L_3 \mid Y)$ (calibration)	$c_{L_3} = 5$
0d	$b < 0.5 \Rightarrow \text{regenerate}$	(decision rule, not a stat)	—
1a	Project b from step 0 to step 1	Kernel (from iter trajectories)	0
1b	Compute expected refinement gain	Kernel + reward	0
1c	Refine + repeat from 0a	(loop back to single-shot stats)	$c_g + \dots$

Key insight: the controller smoothly mixes single-shot estimates (step 0 within an attempt) with the kernel (across attempts). Same Bayesian machinery throughout, just different parameters at different timescales.

The five-benchmark cube

Benchmark	Pool	n inst.	Coverage	k
LCB-hard	102	102	100%	3
LCB-medium	207	207	100%	3
LCB-easy	135	135	100%	3
SWE-bench Lite	300	$\sim 29-47$	17%	3
SWE-bench Verified	500	$\sim 27-30$	6%	3
MBPP+	378	100	26%	3
HumanEval+	164	100	61%	3

Per (gen, bench) cell: n instances $\times k=3$ patches/instance $\Rightarrow \sim 90-620$ records. **LCB pool**: union of all 6 LCB releases (test.jsonl + test2-test6), filtered to leetcode platform $\Rightarrow 444$ leetcode problems total (102 hard + 207 medium + 135 easy). **We sample 100% of the leetcode pool.** **Why SWE varies per generator:** repo-level diffs require multi-file edits; weaker / open-weight gens fail patch-parse more often. Same 50 (Lite) / 30 (Verified) candidates per gen; what differs is how many produce a parseable patch.

The five-benchmark cube — why each + generators

Why each benchmark:

- ▶ **LCB (LiveCodeBench)** has L_2 (public tests) — the only benchmark in our cube with all 4 critics (L_0 syntax, L_1 lint, L_2 public tests, L_3 LLM review). Type: *competitive programming, leetcode-style, function-level*.
- ▶ **SWE-bench (Lite + Verified)** has a fine-grained oracle (FAIL_TO_PASS + PASS_TO_PASS test sets). Type: *repo-level patches, multi-file diff format*.
- ▶ **MBPP+ / HumanEval+** probe the saturation regime (priors ≥ 0.7 on closed-API gens) — they exist in the cube precisely to verify Bayesian collapses to `always_verify` when there's nothing to choose. Type: *function-level, inline asserts*.

Generators (5 total):

- ▶ **4 closed-API frontier** (via OpenRouter): `gpt5_mini`, `qwen3_coder`, `haiku45`, `sonnet45`.
- ▶ **1 open-weight, locally hosted**: `Qwen2.5-Coder-32B-Instruct` via `vLLM` at `:8003` on `Runpod-1` (no API spend, replicable).

LiveCodeBench — sample task

Difficulty: **hard** (LeetCode contest level)

Title: *Maximum Length of Subarray With Limited Frequencies*

Statement (excerpt):

You are given an integer array `nums` and an integer `k`. The frequency of an element `x` is the number of times it occurs in an array.

An array is called `good` if the frequency of each element in this array is less than or equal to `k`.

Return the length of the longest `good` subarray of `nums`.

Public tests (LCB has these built in):

```
assert longestGoodSubarray([1,2,3,1,2,3,1,2], 2) == 6
assert longestGoodSubarray([1,2,1,2,1,2,1,2], 1) == 2
```

Per instance: we sample 3 attempts (`patch_id` $\in$ $\{0, 1, 2\}$).

SWE-bench Verified — sample task

Repository-level bug, real GitHub issue.

Instance: django__django-12284

Issue: *“Models inheriting from a concrete abstract model with a foreign key in the abstract model fail to initialise correctly.”*

Spec hands the LLM:

- ▶ the issue text
- ▶ the buggy commit hash
- ▶ “oracle files” — the files the gold patch touches (read-only context)
- ▶ FAIL_TO_PASS test names **(must pass after fix)**
- ▶ PASS_TO_PASS test names **(must stay passing)**

LLM proposes a unified diff. Harness:

1. Spins up podman container with the buggy commit + correct deps
2. Applies patch
3. Runs FAIL_TO_PASS + PASS_TO_PASS test sets
4. Records $Y=1$ iff **all FAIL_TO_PASS pass** and **no PASS_TO_PASS regress**.

Sampling and seeds

- ▶ **Reproducibility:** seed=42 fixes both the instance sample AND the per-attempt generation across all generators. Same instances and same prompt for every model.
- ▶ **Per-instance attempts:** 3 calibration patches ($\text{patch_id} \in \{0, 1, 2\}$) + up to 4 iter-refine steps.
- ▶ **Total per cell:** $30 \text{ instances} \times 3 + 30 \times 4 = 210$ patch evaluations.
- ▶ **Total cells:** $5 \text{ generators} \times \{\text{LCB} \times 3, \text{MBPP+}, \text{HumanEval+}, \text{SWE-Lite}, \text{SWE-Verified}\} = 5 \times 7 = \mathbf{35 \text{ calibration cells}}$.
- ▶ **Total iter cells:** $5 \text{ generators} \times \{\text{LCB} \times 3, \text{SWE-Lite}, \text{SWE-Verified}\} \times 2 \text{ methods (Self-Refine, Reflexion)} = \mathbf{50 \text{ iter cells}}$.
- ▶ **Aggregate:** $\sim \mathbf{2400 \text{ LLM patches}}$, $\sim \mathbf{500 \text{ harness evaluations}}$.

Calibration phase — what data we collect

For each (generator, benchmark, difficulty, instance, patch_id):

1. Generate patch via OpenRouter / vLLM $\rightarrow$ predictions_p{0,1,2}.jsonl
2. Run cheap critics:
 - ▶ L_0 : ast.parse on each modified .py file (free)
 - ▶ L_1 : ruff check on each modified file (free)
 - ▶ L_2 : public-test runner (LCB-only, free)
 - ▶ L_3 : Haiku-4.5 PASS/FAIL judge ($\sim$ \$0.005/patch, OpenRouter)
3. Run harness oracle:
 - ▶ LCB: hidden-test runner (free)
 - ▶ SWE: **podman container + test suite** (slow but free)
 - ▶ MBPP+/HumanEval+: assert-based unit tests (fast, free)
4. Persist as critic_results.jsonl with (L_0, L_1, L_2, L_3, Y) per patch.

Total budget: $\sim$ \$15 across all 35 calibration cells.

Iterative refinement phase — what data we collect

Real implementations of Self-Refine and Reflexion (commit c5b0faa):

- ▶ Per-instance trajectory: $\text{step}_0 \rightarrow \text{step}_1 \rightarrow \dots \rightarrow \text{step}_4$ with method-specific stop condition.
- ▶ **Comprehensive logging:**
 - ▶ `iter_records.jsonl` — one row per (instance, step) with diff, L_0, L_1, L_3, Y , stop-decision, cost
 - ▶ `raw_calls/` — full LLM call logs (prompts, responses, tokens, cost)
 - ▶ `transition_kernel.json` — measured $P(Y_{k+1} \mid Y_k)$ per cell
 - ▶ `policy_comparison.json` — Bayesian-DP vs Greedy vs threshold vs fixed
 - ▶ `stop_distribution.json` — at which step did the trajectory stop
 - ▶ `RUN_CONFIG.json` — temperature, prompt-template hashes, git SHA

Why this much logging: reproducibility under reviewer pressure. We must be able to answer “how was this Y computed?” for any single record.

One iter_records row — example

```
{
  "step": 1, "instance_id": "django__django-12284",
  "method": "selfrefine", "diff": "--- a/django/db/...",
  "Y": 1,
  "L0_syntax": true,
  "L1_lint": true,
  "L3_llm_review": true,
  "step_cost_usd": 0.012,
  "method_specific": {
    "critique_text": "...the patch correctly handles the abstract..."
  }
}
```

Each row records the FULL state at that step, including:

- ▶ the patch text (so we can re-evaluate later if critics change)
- ▶ all critic outcomes
- ▶ ground-truth Y from the harness
- ▶ the LLM's self-critique or reflection text
- ▶ monetary cost

Tier-1 statistical rigor (commit 54218f6)

Five techniques, all computed before claiming any positive result:

1. **Paired-bootstrap CIs** ($B = 1000$) on per-instance $\Delta = U_\pi - U_{\text{always_verify}}$.
2. **χ^2 test of critic independence** — checks that $L_0 \perp L_1 \perp L_3 \mid Y$ assumption holds.
3. **MDE-power** — for each cell, smallest detectable effect at $\alpha = 0.05$, $\beta = 0.2$. We have $n = 29$ on LCB-hard $\Rightarrow$ MDE ≈ 7 utility units.
4. **LOO-CV** — leave-one-instance-out cross-validation. Train likelihoods on $n - 1$ instances, evaluate the controller on the held-out one. Repeat n times.
5. **Benjamini-Hochberg FDR** per policy at $q = 0.05$ across the 28-cell cube.
20 of 28 Bayesian-greedy cells survive BH-FDR.

LOO-CV resolves the train/test concern

Concern raised: likelihoods and trajectory data are computed on the *same* instances we evaluate on. Reviewer would ask: **are policy comparisons overfitting?**

Resolution (commit 5ae5c64):

- ▶ Run leave-one-out CV on all 12 LCB cells.
- ▶ For each held-out instance, compute likelihoods on the remaining $n - 1$ patches, then evaluate every policy on the held-out instance.
- ▶ Aggregate: per-instance Bayesian win rate vs. in-sample Bayesian win rate.

Result: differences are within ± 1 utility unit on average — the in-sample numbers are honest.

Why it works: likelihoods are estimated with Beta(1,1) smoothing, so 1-instance leave-out shifts $P(z \mid Y)$ by at most $\sim 2\text{-}3\%$. The controller's decisions are robust to that.

Cluster bootstrap on SWE iter cells

Why: SWE iter trajectories are auto-correlated within an instance (step k 's diff conditions on step $k - 1$). Standard per-row bootstrap underestimates SE.

Cluster bootstrap (commit 54218f6):

- ▶ Resample at the *instance* level (with replacement).
- ▶ Re-compute trajectory utility per resample.
- ▶ Compute bootstrap CI from the empirical distribution.

Effect on SE estimates:

- ▶ Per-row bootstrap SE: σ_{row}
- ▶ Cluster bootstrap SE: $\sigma_{\text{cluster}} \approx 1.6 \cdot \sigma_{\text{row}}$ (**rows are auto-correlated within instance**)
- ▶ CIs widen by $\sim 60\%$ but **the headline $\Delta > 0$ still survives** on most cells.

Pre-registered prediction (commit 34713d8)

Before running the LCB experiment, we wrote a prediction in `PRE_REGISTRATION.md`:

Pre-registered claim (S4)

“On LCB-hard with gpt5_mini, the Bayesian-greedy controller with the calibrated likelihood table will yield $\Delta U \in [+8, +18]$ utility units vs. `always_verify`, paired-bootstrap 95% CI strictly above 0.”

Empirical result: $\Delta = +12.55$, **CI** $[+5.5, +19.5]$.

Synthetic forecast on calibration data alone: $+12.2$. **Why this is rare:** most

empirical ML papers don't pre-register predictions. This is direct evidence the framework's predictive structure generalises.

Notation hint — read once, applies to all result slides

Outcomes & cells

- ▶ $Y \in \{0, 1\}$ — ground-truth correctness of a generated patch ($Y=1$ passes the full hidden test suite; $Y=0$ broken).
- ▶ $z \in \{\text{pass}, \text{fail}\}$ — a critic's verdict on a patch (cheap proxy for Y).
- ▶ Cell = (generator, benchmark) pair. We have $5 \text{ gens} \times 7 \text{ benches} = 35 \text{ cells}$.

Critic stack (in order of cost)

- ▶ L_0 — syntax / lint / static parse (cheapest).
- ▶ L_2 — public test execution / fast partial tests.
- ▶ L_3 — LLM-judge review of the patch.
- ▶ verify — full hidden test suite (gold-standard label for Y ; expensive).

(No L_1 in our system: we skip from L_0 to L_2 to reflect the natural cost gap between syntactic checks and test execution.)

Probabilities & informativeness

- ▶ $prior = P(Y=1)$ — generator's first-shot success rate on this cell, measured from the calibration corpus.
- ▶ $P(\text{pass} \mid Y=1)$ — critic recall (true-positive rate).
- ▶ $P(\text{pass} \mid Y=0)$ — critic false-positive rate.
- ▶ $L_2 \text{ gap (informativeness)}$
 $= P(\text{pass} \mid Y=1) - P(\text{pass} \mid Y=0)$.
 $0 = \text{noise}$; $1 = \text{oracle}$

Cost model & utility

- ▶ Default cost vector: $c_{\text{gen}}=5$, $c_{L_0}=1$, $c_{L_2}=2$, $c_{L_3}=5$, $c_{\text{ver}}=30$, reward $R=100$.
- ▶ Key ratios: $c_{\text{ver}}/c_{L_2}=15$, $R/c_{\text{ver}} \approx 3.3$.
- ▶ Utility per query: $U = R \cdot \mathbf{1}[\text{final patch correct}] - \sum (\text{action costs paid})$.
- ▶ $\Delta_{\text{vs always_verify}} = U(\text{policy}) - U(\text{always_verify})$, mean over instances; reported per cell.

Policies compared

- ▶ bayesian_greedy — 1-step lookahead: at each step, pick the action that maximises immediate expected utility under the current belief $P(Y=1 \mid \text{evidence so far})$.
- ▶ bayesian_DP — full Bellman recursion over (belief, budget) state. Multi-step optimal under the model; near-identical Δ to greedy in all our cells.
- ▶ threshold_L0 / L2 / L3 — run that one critic; verify only if it says “fail”; otherwise accept.
- ▶ best_of_3 — generate 3 patches, verify each, return the best.
- ▶ fixed_pipeline — hard-coded sequence $L_0 \rightarrow L_2 \rightarrow \text{verify}$.
- ▶ always_verify — skip critics, always pay c_{ver}

Three regimes — the paper’s organising claim

Notation: $prior = P(Y=1)$, the probability the generator’s first single-shot patch is correct (measured per cell from the calibration corpus). $L_2 \text{ gap} = P(\text{pass} \mid Y=1) - P(\text{pass} \mid Y=0)$, the informativeness of the mid-tier critic.


Bayesian variant reported: *bayesian_greedy* (1-step lookahead). *bayesian_DP* (full Bellman recursion) gives near-identical Δ in all three regimes — both are listed as “Best policy” below.

Cost vector for the policy comparison: $c_{\text{gen}}=5$, $c_{L_0}=1$, $c_{L_2}=2$, $c_{L_3}=5$, $c_{\text{ver}}=30$, $R=100$. Key ratios: $c_{\text{ver}}/c_{L_2}=15$, $R/c_{\text{ver}} \approx 3.3$. Sensitivity to these choices is on the 4×4 (R, c_{ver}) grid slide.

Regime	Description	Defining condition	Best policy
A	low prior, informative cheap critics	LCB-hard, prior ≈ 0.2 , $L_2 \text{ gap} \approx 0.49$	Bayesian DP / greedy
B	mid prior, near-oracle L_2 available	SWE-Lite / Verified, prior ≈ 0.4	Bayesian DP / greedy
C	high prior, saturated benchmark	MBPP+ / HumanEval+ (prior > 0.9), LCB-easy on strong gens	threshold_L2 ties always_verify

Defensible framing: we don’t claim “Bayesian wins everywhere.” We claim a regime characterisation. Reviewers can’t refute it with a single counter-example because the regime conditions are explicit.

Three-regime map — visual



x-axis: prior $P(Y=1)$ (generator strength on benchmark). y-axis: L_2 gap $P(\text{pass} | Y=1) - P(\text{pass} | Y=0)$.

Color: Bayesian-greedy Δ utility vs always-verify at the default cost vector ($c_{\text{gen}}=5$, $c_{L_0}=1$, $c_{L_2}=2$, $c_{L_3}=5$, $c_{\text{ver}}=30$, $R=100$); green = Bayesian wins, yellow ≈ 0 tie. Robustness to cost ratio: see the 4×4 (R , c_{ver}) grid two slides later.

Per-cell winner — benchmark × generator matrix

Cost vector held fixed at default deployment regime: $R=100$, $c_{\text{ver}}=30$, $c_{L_0}=1$, $c_{L_2}=2$, $c_{L_3}=5$, $c_{\text{gen}}=5$.

Ratios: $c_{\text{ver}}/c_{L_2}=15$, $R/c_{\text{ver}}\approx 3.3$. Per-cell winner shifts at other ratios (see 4×4 grid two slides later).

LCB-hard	bayes_g $\Delta = +13.0$	bayes_DP $\Delta = +20.5$	bayes_DP $\Delta = +23.1$	bayes_DP $\Delta = +20.2$	bayes_DP $\Delta = +18.2$
LCB-medium	thr(L2) $\Delta = +14.6$	bayes_g $\Delta = +18.7$	bayes_g $\Delta = +19.9$	bayes_g $\Delta = +17.5$	bayes_DP $\Delta = +6.5$
LCB-easy	thr(L2) $\Delta = +21.4$	thr(L2) $\Delta = +9.7$	bayes_g $\Delta = +17.3$	thr(L2) $\Delta = +11.7$	thr(L2) $\Delta = +1.2$
MBPP+	always_v $\Delta = +0.0$	thr(L2) $\Delta = +0.6$	thr(L2) $\Delta = +0.3$	always_v $\Delta = +0.0$	thr(L2) $\Delta = +1.9$
HumanEval+	thr(L2) $\Delta = +1.4$	thr(L2) $\Delta = +4.2$	always_v $\Delta = +0.0$	always_v $\Delta = +0.0$	thr(L2) $\Delta = +2.4$
SWE-Lite	thr(L2) $\Delta = +13.1$	thr(L2) $\Delta = +11.8$	thr(L2) $\Delta = +15.5$	thr(L2) $\Delta = +21.9$	bayes_g $\Delta = +9.3$
SWE-Verified	thr(L2) $\Delta = +18.6$	thr(L2) $\Delta = +4.7$	thr(L2) $\Delta = +13.8$	thr(L2) $\Delta = +9.3$	bayes_g $\Delta = +1.7$

Per-cell winner
■ bayes_g
■ bayes_DP
■ thr(L2)
■ thr(L0)
■ thr(L3)
■ always_v

gpt-5-mini (API)
qwen3-coder (API)
haiku-4.5 (API)
sonnet-4.5 (API)
qwen2.5-32b (local)

Each cell: winning policy (bold) + its Δ vs always_verify. **Green** = regime A (Bayesian wins); **orange** = 51 / 87

Per-cell winner — counts & regime framing

Winner	n / 35	Regime / typical cells
bayesian_greedy	10	Regime A — LCB-hard / -medium / -easy on closed-API gens.
bayesian_DP	2	Regime A (within ϵ of greedy): LCB-hard / haiku, qwen32b.
threshold(L2)	17	Regime B — all SWE-Lite & SWE-Verified closed-API; many MBPP+/HumanEval+.
always_verify	4	Regime C ties: MBPP+ {gpt5, sonnet}; HumanEval+ {haiku, sonnet}.
threshold(L0/L3)	2	Open-weight SWE quirks: SWE-Lite/qwen32b (L3); SWE-Ver/qwen32b (L0).


Honest framing: Bayesian (greedy + DP) wins 12/35 cells — regime A, where critics are noisy and need composing. threshold(L2) wins 17/35 — regime B, where the public-test critic has gap > 0.7 (just trust the tests). always_verify wins/ties 4/35 regime-C saturated cells.

Paper claim should be regime-conditional: “Bayesian dominates regime A” *not* “Bayesian wins”.

Per-cell winner — 4×4 grid varying R (rows) and c_{ver} (cols)






Critic costs held fixed at default per-critic values: $c_{L_0}=1$, $c_{L_2}=2$, $c_{L_3}=5$, $c_{\text{gen}}=5$. Only R and c_{ver} vary; the ($R=100$, $c_{\text{ver}}=30$) panel reproduces slide 51 cell-by-cell.

Winner per (benchmark, generator) under varying R and c_{ver} — 7 benches $\times$ 5 gens per panel



Each panel = one (R , c_{ver}). Across 560 (panel $\times$ cell) outcomes at default critic costs: always_v 33%, bayesian_DP 30%, thr(L2) 29%, bayes_g 8%, thr(L0) 1%, thr(L3) 0%.






All policies, per model — LCB-hard



Take-aways

★ = winner per cell. **bayes_greedy / DP** win by +12 to +20 on the 4 closed-API gens. **qwen2.5-32b** only +4.6, CI crosses 0. **Self-Refine, Reflexion**: -12 to -27 across all gens.






All policies, per model — LCB-medium



Take-aways

bayes_greedy wins on 4 of 5 gens (+15 to +19). **qwen2.5-32b** is the **regime crossover**: threshold(L2) wins +2.4; bayes loses at -5.9. **Self-Refine** ~ -10 to -20 , Reflexion similar.






All policies, per model — LCB-easy



Take-aways

gwen3 / haiku / sonnet: bayes_greedy +19–+20. **gpt-5-mini:** threshold(L2) +20.9 and Reflexion +38.4 beat bayes (+5.5, +18.0). **qwen2.5-32b:** all near 0 – saturated prior 0.97.






All policies, per model — MBPP+



Take-aways

Regime C: 4 closed-API gens cluster at $\Delta \approx 0$ for bayes / threshold(L2). **qwen2.5-32b:** threshold(L2) +1.9, bayes -2.0 – threshold wins. **All other policies negative** (refinement / extra critics waste budget).





All policies, per model — HumanEval+



Take-aways

Regime C strongest here (prior 0.84–0.94). **threshold(L2) wins on 3 of 5 gens**: gpt5 +1.4, qwen3 +4.2, qwen32b +2.4. bayes_greedy / DP ≈ 0 (correctly defaults to always_verify). Big losses for L3 / fixed_pipeline (–12 to –20).





All policies, per model — SWE-Lite



Take-aways (SWE-Lite, all 4 closed-API gens)

threshold(L2) wins on every gen: gpt5 +13.1, qwen3 +11.8, haiku +15.5, sonnet +21.9. bayes_DP is ~2–3 utility points behind threshold(L2) but still beats greedy by +5 on haiku. **Reflexion catastrophic** (–67 to –72); Self-Refine –13 to –41.

All policies, per model — SWE-Verified




Take-aways (SWE-Verified, all 4 closed-API gens)

threshold(L2) wins on every gen: gpt5 +18.6, qwen3 +4.7, haiku +13.8, sonnet +9.3. bayes_DP closer to threshold(L2) than greedy (+10.7 vs +4.2 on haiku). **fixed_pipeline catastrophic** (−55 to −69); Self-Refine −16 to −30.

Sensitivity 1/3: c_{ver} sweep on LCB-hard — 5 generators

c_{ver} sweep on LCB-hard --- regime shifts as verifier cost grows



Sweep $c_{\text{ver}} \in \{10, 15, 20, 30, 40, 60, 100\}$ keeping reward $R=100$ and likelihoods fixed; $n=29$ instances per generator. Each panel: utility curves for 8 policies. Next slide: regime interpretation.

Sensitivity 2/3: cost regimes & implication for the headline

Default cost vector used in the headline figures

$R = 100$, $c_{\text{ver}} = 1$, $c_{L_0} = c_{L_2} = c_{L_3} = 0.05$, $c_{\text{gen}} = 1 \Rightarrow \text{ratio } c_{\text{ver}}/R = 0.01$
(verification very cheap).






Three cost-regimes visible on every generator:

1. $c_{\text{ver}}/R \lesssim 0.15 \rightarrow \text{always_verify} / \text{threshold}(L_2)$ wins. Verification cheap, oracle on every patch beats anything fancier.
2. $0.2 \lesssim c_{\text{ver}}/R \lesssim 0.5 \rightarrow \text{Bayesian dominates}$: its give-up-early ability pays for the critic chain.
3. $c_{\text{ver}}/R \gtrsim 0.6 \rightarrow \text{all policies collapse to 0}$ (always give up). Verifying too expensive for any policy.

Implication for the paper:

- ▶ Headline regime ($c_{\text{ver}}/R = 0.01$) is regime (i) — the *least favourable* for Bayesian.
- ▶ Bayesian's win there comes from *accuracy on hard instances*, not cost-arbitrage.
- ▶ Deployers with expensive verifiers (LLM-judge, slow CI) enter regime (ii); advantage *grows*.
- ▶ **Action item**: report the c_{ver}/R ratio explicitly in the abstract.






c_{ver} snapshot — Regime-A cell (LCB-hard / gpt-5-mini)



What changes as c_{ver} grows (LCB-hard / gpt-5-mini, $R=100$)

$c_{ver}=10$: **threshold(L2)** wins (+21.2); always_verify close behind. $c_{ver}=20$: **bayes_greedy** wins (+18.4) — crossover. $c_{ver}=30$: **bayes_greedy** (+13.6) clearly best; threshold(L2) drops to +7.4. $c_{ver}=60, 100$: **collapse regime** — only Bayesian survives at 0; every other policy goes deeply negative. **Star** marks the per-cell winner.






Reward snapshot: same cell, 5 different reward magnitudes



What changes as reward R grows (LCB-hard / gpt-5-mini, $c_{\text{ver}}=30$)

$R=25, 50$: reward not worth the cost; everything negative; **Bayesian gives up at 0** (the only winner).
 $R=100$: **bayes_greedy wins** (+13.6); the regime where Bayesian's edge is largest. $R=200, 500$: reward dominates costs; **threshold(L2) overtakes** (+41.9, +145.3) — when verifying is essentially free relative to reward, just trust L2. Bayesian DP catches up but greedy lags. **Insight**: the c_{ver}/R ratio drives the regime — doubling R is equivalent to halving costs.


c_{ver} snapshot — Regime-B cell (SWE-Verified / sonnet-4.5)



What changes as c_{ver} grows (SWE-Verified / sonnet-4.5, $R=100$)

threshold(L2) wins for $c_{\text{ver}} \leq 60$ (gap +5 to +8 over Bayesian) — L_2 public-test critic is near-oracle on SWE so trusting it dominates. At $c_{\text{ver}}=100$, threshold(L2) collapses to -7.7 (still verifies often, pays a fortune); **Bayesian gives up** and lands at 0 — $+7.7$ ahead. **This cell is why the headline “Bayesian wins” is regime-conditional:** when L_2 is near-oracle, just trust the tests.

Sensitivity 3/3: θ -perturbation — does miscalibration kill the win?



Take-aways

Test: for each $P(z|Y)$ entry in `likelihood_tables.json`, perturb by $\pm 10\%$ / $\pm 20\%$, refit both Bayesian controllers, re-run all 8 policies on the same trajectories. alt $\pm 20\%$ = entries perturbed in alternating directions (adversarial worst-case).

Result: Bayesian Δ moves by at most ~ 2 utility points under $\pm 20\%$ perturbation. The 20% bar crosses the clean bar's CI on every closed-API gen.

qwen2.5-32b is the exception: CIs are wider and one perturbation (+20%) drops Δ from +4.6 to +2.4. Consistent with the 3-identical-patches behaviour reported in §Honest negatives.

Conclusion: headline result is robust to realistic likelihood miscalibration; we are not over-fitting on a knife-edge calibration.

Bayesian wins after BH-FDR correction

Policy	cells	uncorr. $p < .05$	per-policy BH	global BH
bayesian_greedy	28	20	20	18
bayesian_DP	28	22	21	19
threshold_L2	28	21	20	18
threshold_L3	28	26	25	23
threshold_L0	28	24	24	22
fixed_pipeline	28	28	28	27
best_of_3	28	24	24	22

Bayesian wins are stable under multiple-comparisons correction.

3 borderline cells (gpt5/qwen3/sonnet on SWE-Verified at $\Delta \in [+3, +4]$) sit on the BH boundary; reported as such.

fixed_pipeline appears strongest because it always loses to `always_verify` ($\Delta < 0$ everywhere) — 28/28 “significant” means **uniformly worse**.

Where Bayesian wins — top 10 cells by Δ utility

$\Delta = U(\text{bayesian_greedy}) - U(\text{always_verify})$ per query. *bayesian_DP* gives near-identical Δ . Cost vector: $c_{\text{gen}}=5$, $c_{L_0}=1$, $c_{L_2}=2$, $c_{L_3}=5$, $c_{\text{ver}}=30$, $R=100$ ($c_{\text{ver}}/c_{L_2}=15$, $R/c_{\text{ver}}\approx 3.3$); robustness on the 4×4 grid slide.

Generator	Benchmark	Regime	Δ (bayes_greedy – always_verify)
sonnet45	LCB-hard	A	+20.3
qwen3_coder	LCB-hard	A	+20.3
haiku45	LCB-hard	A	+20.3
haiku45	LCB-easy	A/B	+20.3
qwen3_coder	LCB-easy	A/B	+19.8
sonnet45	LCB-easy	A/B	+19.3
haiku45	LCB-medium	A	+19.1
qwen3_coder	LCB-medium	A	+18.4
sonnet45	SWE-Lite	B	+17.6
sonnet45	LCB-medium	A	+17.3

Pattern: biggest wins all in regime A (low prior, high L_2 gap). Regime B (SWE) wins are smaller but still uniformly positive.

Common factor: every win in this list has L_2 gap $> +0.40$ (highly informative critic) and prior < 0.55 (uncertain enough that gating helps).

Case study: sonnet45 $\times$ LCB-hard ($\Delta = +20.3$)

Cell parameters:

- ▶ Prior $P(Y=1) = 0.18$ (capable model, but LCB-hard is hard)
- ▶ L_2 gap = +0.49 (workhorse signal: 90% of correct patches pass public tests; 41% of broken ones do too)
- ▶ L_3 gap = +0.07 (mildly informative)

Cost analysis per instance:

Strategy	Cost	Outcome
always_verify	$5 + 30 = 35$	always pay c_v to learn Y
bayesian_greedy	$5+5$ if reject, $5+5+30$ if accept	gate via L_2 , only verify confident ones

Aggregate per instance ($n = 30$ cohort):

- ▶ always_verify: $\bar{U} = 100 \cdot 0.18 - 35 = -17.0$
- ▶ bayesian_greedy: $\bar{U} = +3.3$
 - ▶ L_2 fails on ~ 0.59 of patches $\rightarrow$ skip verify, save $30 \times 30 \times 0.59 = \531 across cohort
 - ▶ L_2 passes on $\sim 0.41 \rightarrow$ verify, capture reward
- ▶ $\Delta = +20.3$ utility per instance, paired-bootstrap CI $[+13.7, +25.6]$

Per-instance: when Bayesian saves money

Concrete instance: LCB-hard #3203, gpt5_mini patch p_0 .

```
diff: ~800 chars (a sliding-window attempt)
L0_syntax: True      (Python parses)
L2_public_tests: False (1 of 3 public examples FAILS)
L3_llm_review: False (judge says "looks wrong")
Y (oracle): 0        (this patch doesn't fix the bug)
```

Two policies side-by-side:

`always_verify`

- ▶ Pay $c_g + c_v = 5 + 30 = 35$
- ▶ Get reward 0 ($Y=0$)
- ▶ Utility: **-35**

`bayesian_greedy`

- ▶ Pay $c_g + c_{L_2} = 5 + 5 = 10$
- ▶ L_2 fails $\rightarrow$ posterior drops well below 0.5 $\rightarrow$ skip verify
- ▶ Utility: **-10**

Per-instance saving: +25 utility units. And we didn't lose anything — the patch was actually $Y=0$, so verifying would only have confirmed it cost-fully.

Per-instance: when Bayesian correctly verifies

Concrete instance: LCB-hard #2784, qwen3_coder patch p_1 .

```
L0_syntax: True
L2_public_tests: True    (all 3 examples PASS)
L3_llm_review:   True    (judge: "looks correct")
Y (oracle):      1       (resolves the bug!)
```

Two policies side-by-side:

`always_verify`

- ▶ Pay $c_g + c_v = 5 + 30 = 35$
- ▶ Get reward 100
- ▶ Utility: **+65**

`bayesian_greedy`

- ▶ Pay $c_g + c_{L_2} + c_v = 5 + 5 + 30 = 40$
- ▶ L_2 passes $\rightarrow$ posterior climbs to $\sim 0.65 \rightarrow$ verify
- ▶ Get reward 100 (correctly accepted)
- ▶ Utility: **+60**

Bayesian pays 5 extra (c_{L_2}). But this 5-unit “overhead” is the price of the gating that saves 25 on the OTHER $\sim 59\%$ of instances (per the case-study slide).
Net per-cohort win: +20.3.

Bayesian doesn't beat `always_verify` on every instance — it wins on average by gating cheaply and verifying selectively.

LCB iter cells — replay vs. real implementation

Important honest finding: our policy-replay approximations of Self-Refine and Reflexion diverged from real implementations on 18 of 24 cells, sometimes **flipping the sign**.

Difficulty	Generator	Self-Refine real (replay)	Reflexion real (replay)
LCB-hard	gpt5_mini	−33.5 (−26.9)	−5.2 (−11.4)
LCB-hard	qwen3_coder	−11.4 (−23.5)	−21.7 (−22.8)
LCB-medium	gpt5_mini	−1.7 (−10.0)	+4.3 (+1.8)
LCB-easy	gpt5_mini	+3.0 (+16.7)	+27.7 (+38.4)
LCB-easy	haiku45	−0.7 (−22.0)	+27.3 (−24.4)

(Difference: replay → real, on a few telling cells.)

Why the divergence: replay assumed monotonic improvement. Real Self-Refine often *declines* after early CRITIQUE_OK. Real Reflexion sometimes finds a fix on step 1 that replay would assign to step 4.

Today's claim: Bayesian framework wins 22/24 LCB cells under real implementation; 2 Reflexion-wins cells reported transparently.

Open-weight integration: Qwen2.5-Coder-32B

Why: biggest reviewer pushback would be “where’s the open-weight generator?” We added qwen25_32b as the **5th generator panel column** — local vLLM serving Qwen2.5-Coder-32B-Instruct, no API spend.


Calibration headline (LCB):

Difficulty	n	prior	L_2 gap	Bayesian-DP Δ
hard	29	0.21	+0.49	+4.62 [−3.5, +11.7]
medium	90	0.36	+0.28	−7.83 [−10.8, −4.2]
easy	62	0.97	+0.85	−0.69 [−2.6, +2.5]

Pattern: Qwen32B’s open-weight column **reinforces the regime story**. On hard (regime A) Bayesian wins; on medium it loses; on easy nothing matters. Same as closed-API panel.

SWE-bench (calibration done): SWE-Lite: $n=82$, prior 0.21, L_3 gap 0.44 $\Rightarrow$ **thr(L3) wins +8.79** (regime-B-ish on the open-weight panel). SWE-Verified: $n=77$, prior 0.30, L_3 gap 0.55 $\Rightarrow$ **thr(L0) wins +4.56**. Per-cell winners visible in the matrix slide.

L3-reviewer invariance — Bayesian's edge is robust



L3 informativeness gap on LCB-hard. Bold border = self-review.

gpt5_mini and **qwen3_coder** show clean reviewer-invariance. **sonnet45/sonnet45** self-review collapses to negative gap (the model can't critique itself reliably). We use **haiku45** as the default reviewer across the panel.

Honest negative findings (we report them)

1. mypy critic was a no-op

- ▶ Adding mypy as a 5th cheap critic showed near-zero gap. Likelihood ratio ≈ 1 at $Y=1$ vs $Y=0$. **Reported as null finding** (commit 0babe59).

2. SWE-bench Pro deferred

- ▶ Pro prompts are $\sim 24\times$ larger than Lite (multi-file oracle context, $\sim 121K$ tokens). gpt-5-mini exhausts `max_tokens` budget on reasoning; produces empty parseable output. Documented in `data/swebench_pro_probe/`; **deferred to follow-up**.

3. LCB iterative refinement HURTS in some regimes

- ▶ Self-Refine on LCB-hard with gpt5-mini: $\Delta = -33$. The model's self-critique is noisy and pushes it to overcorrect.
- ▶ Same model on LCB-easy: $+27$ (Reflexion). Refinement value is benchmark-dependent, not universal.

Permanent harness-side data limitations

~6 instances permanently fail to evaluate:

- ▶ `sympy__sympy-20916`, `sphinx-doc__sphinx-*`, `scikit-learn__scikit-learn-*`, `pydata__xarray-*`
- ▶ **Root cause:** pip dep chains rotted since 2023. `setup_env.sh` fails with exit status 2 — *not* a bug in our pipeline.
- ▶ Same instances would error for any 2026 user of SWE-bench.

How we handle: document explicitly in `HARNESS_ERRORS_NOTE.json`. ~6 data points lost out of ~600 SWE evaluations. **Below-noise.**

Methodology rigor — what we did vs typical findings paper

Technique	Done?	Where
Pre-registration	✓	PRE_REGISTRATION.md
Paired-bootstrap CI	✓	all 28 cells, $B = 1000$
χ^2 critic-independence test	✓	critic_independence/
MDE-power analysis	✓	mde_power/
LOO-CV	✓	all 12 LCB cells
Benjamini-Hochberg FDR	✓	per-policy + global
Wilson CI on priors	✓	prior_ci/
θ sensitivity sweep	✓	lcb_sensitivity.py
c_v sensitivity sweep	✓	cver_sensitivity_sweep.py
Cluster bootstrap on SWE iter	✓	cluster_bootstrap_swe.py
Multi-L3-reviewer invariance	✓	3 reviewers $\times$ 4 gens
Real Self-Refine + Reflexion impls	✓	not just policy replays

Most findings-style papers in this area have $\sim 3-4$ of these.

What's done vs running vs pending

Done ✓

- ▶ LCB calibration: 5 generators $\times$ 3 difficulties = **15 cells**
- ▶ LCB iter (Self-Refine + Reflexion, real impls): **30 cells**
- ▶ MBPP+/HumanEval+ calibration: **8 cells**
- ▶ SWE-bench Lite + Verified calibration: **8 cells** (4-gen, original data)
- ▶ Tier-1 methodology rigor batch (FDR, MDE-power, LOO-CV, etc.)
- ▶ All 4 swebench podman patches (patches/swebench/)
- ▶ Qwen32B SWE calibration: **4 of 6 reports**

Running 🔄

- ▶ Artem-1: verified/qwen25_32b/selfrefine step 3 in flight
- ▶ Artem-3: verified/qwen25_32b/reflexion step 1 in flight

Pending ○

- ▶ Y-backfill across SWE iter cells
- ▶ Compute kernel + policy_comparison for all SWE cells
- ▶ Refresh PAPER_TABLE + figures (5 generators $\times$ all cells)
- ▶ Paper submission draft

Final cell count

Region	Cells	Status
LCB calibration (5 gen $\times$ 3 diff)	15	✓ done
LCB iter (5 gen $\times$ 3 diff $\times$ 2 method)	30	✓ done
MBPP+ / HumanEval+ calibration (4 gen $\times$ 2 bench)	8	✓ done
SWE calibration (5 gen $\times$ 2 bench)	10	🕒 8/10, qwen32b in flight
SWE iter (5 gen $\times$ 2 bench $\times$ 2 method)	20	🕒 12/20 done
Total cells with real Y	~83	🕒 ~90%

Compare to typical findings-style paper in this area: 1-3 benchmarks, 2-3 baselines, no pre-registration, no FDR. We're at the high end of empirical depth.

Adding a new benchmark — 13-step pipeline

Required artifacts: HF dataset OR local jsonl, oracle execution mechanism (inline asserts / public+hidden tests / per-instance container).

Steps 1–6 (calibration, IID-synthesized kernel):

1. Register dataset in `scg.sample_instances()`
2. Sample $n = 30$ instances with `seed=42`
3. **Calibration:** `spot_check_generators.py` → `predictions_p{0,1,2}.jsonl`
4. Harness eval → Y per (instance, patch_id)
5. `calibrate_from_spotcheck.py` → `likelihood_tables.json`
6. `lcb_compare.py` → `policy_comparison.json` (IID kernel placeholder)

Steps 7–13 (iter, measured kernel):

7. `iter_refine_real_baselines{,_swe}.py` → `iter_records.jsonl` ($Y=null$)
8. Harness on iter predictions → `eval/*.json`
9. `backfill_swe_iter_y.py` → fills Y
10. ★ `compute_transition_kernel.py` → `transition_kernel.json` (measured)
11. `run_baseline_vs_controller.py --kernel measured`
12. `lcb_summarize_paper.py` → `PAPER_TABLE.json`
13. `lcb_make_figures.py` → refresh figures

Adding a new generator — 13-step pipeline

Required artifacts: model accessible via OpenRouter OR local vLLM.

Setup (steps 1–2):

1. Register in *both* `scg.GENERATORS` and `lcb_calibrate.GENERATORS`
2. For vLLM: ensure two-client pattern in `iter_refine_real_baselines{, _swe}.py`

Calibration phase (steps 3–6, IID kernel):

3. Run `spot_check_generators.py` on every benchmark in the panel
4. Harness eval on calibration patches
5. `calibrate_from_spotcheck.py` → `likelihood_tables.json`
6. `lcb_compare.py` → `policy_comparison.json` (IID)

Iter phase + kernel (steps 7–13):

7. `iter_refine_real_baselines{, _swe}.py`
8. Harness on iter predictions
9. `backfill_swe_iter.y.py`
10. ★ `compute_transition_kernel.py` per (gen, bench, method) → ~6 kernels
11. `run_baseline_vs_controller.py --kernel measured`
12. `lcb_l3_sweep` + `lcb_compare_swap_reviewer` (reviewer-invariance, ~\$0.50)
13. PAPER_TABLE + figures refresh

Cost: ~\$3–5 closed-API; ~\$1 open-weight (only L3 reviewer).

Adding a new critic (1/2) — why the kernel stays

Most invasive change — touches calibration, likelihood schema, Bayesian update, policy comparison.

The kernel is critic-agnostic. Recall the factorization:

$$P(Y_{k+1}, z_\ell \mid Y_k) = \underbrace{P(Y_{k+1} \mid Y_k)}_{\text{kernel}} \cdot \underbrace{P(z_\ell \mid Y_{k+1})}_{\text{critic likelihood}}$$

Critics are *observations of Y*, not transitions. Adding a critic adds a likelihood factor; the kernel is reused unchanged.

Practical consequence: you do **not** need to re-collect iter trajectories or recompute `transition_kernel.json`. Steps below stay within the calibration corpus + likelihood table + Bayesian update; the kernel is a black-box dependency.

Next slide: the 10-step recipe (definition + plumbing + re-evaluation), and the cost.

Adding a new critic (2/2) — 10-step recipe

Steps 1–7 (definition + plumbing):

1. Define `critic_LX(modified_files, problem) -> (PASS, cost)`.
2. Add LX column to `critic_results.jsonl`.
3. Update `likelihood_tables.json` schema with the LX entry.
4. Add pairwise χ^2 tests (LX vs L_0 , L_3).
5. Update `BayesianController`: add likelihood factor to posterior.
6. Add `policy_threshold_LX` for the threshold-policy panel.
7. Add c_{LX} to the cost vector.

Steps 8–10 (re-evaluate, no kernel work):

8. Re-run `calibrate_from_spotcheck.py --skip-generate` → adds LX column without regenerating patches.
9. Re-run `run_baseline_vs_controller.py` (reuses the existing `transition_kernel.json` unchanged).
10. Refresh `PAPER_TABLE.json` + figures.

Cost: \$0 for free critics (lint, mypy, ast); ~\$10–50 for LLM-based critics on the full ~2400-patch calibration corpus.

Master validation checklist (applies to all 3 addition flows)

Check	Threshold	Failure means
Prior $P(Y=1)$	$\in [0.05, 0.95]$	saturation/pathological cell
χ^2 critic-independence	$p > 0.05$	product likelihood unsafe
Critic gap $ \Delta_\ell $	> 0.05 on at least one cell	no-op critic, drop from panel
Reviewer invariance	same Bayesian Δ ranking across 3 reviewers	reviewer-specific artifact
LOO-CV stability	per-instance utility within ± 1 of in-sample mean	likelihoods overfit
Bootstrap CI on Δ	excludes 0 (or borderline reported)	effect indistinguishable from noise
Measured kernel vs IID	differs significantly	F1 ablation pattern as expected

Codified in `PLAYBOOK.md`. New benchmarks/generators/critics extending the panel must pass this checklist before landing in `PAPER_TABLE`.

Genuine weaknesses (we'll surface them)

1. **n is small on headline cells.** 29 LCB-hard, 30 SWE-Verified. **MDE-power passes but optics suffer.** Mitigation: extend LCB-hard to $n \geq 100$ post-submission.
2. **No theoretical contribution.** Empirical-only paper. **Explicitly framed as a findings submission**, not main.
3. **Bayesian gain narrowed under real Self-Refine/Reflexion impl.** Replay said 19/20 wins; real says 22/24 (LCB cells), some Reflexion-wins cells. **Reported transparently — not buried.**
4. **Three regime points still sparse.** SWE-Pro, multilingual code, real \$5+ oracle costs deferred for follow-up work.
5. **Single user-side rate-limit IP for HF.** Caused some retry pain during multi-machine sharding. **Logistical, not research.**

One-page TL;DR

- ▶ **Problem:** when should an LLM-patch pipeline verify vs. regenerate?
- ▶ **Method:** Bayesian POMDP controller using cheap critic likelihoods + measured transition kernel.
- ▶ **Pre-registered prediction confirmed:** synthetic +12.2 → empirical **+12.55** on LCB-hard gpt5_mini.
- ▶ **Three-regime characterisation:** Bayesian dominates regime A (low prior + informative critics), wins regime B (mid prior, oracle available), saturates in regime C (high prior).
- ▶ **Holds under:** BH-FDR, paired bootstrap, LOO-CV, χ^2 critic-independence, MDE-power, real Self-Refine/Reflexion implementations.
- ▶ **Open-weight reproduction:** Qwen2.5-Coder-32B reinforces the regime story, not just frontier-API artifact.
- ▶ **Honest negatives reported:** mypy null, LCB-hard refinement-hurts, SWE-Pro deferred, ~ 6 permanent-error instances.
- ▶ **Submission target:** 3-week window.

Questions / discussion

Repo: `rvz16/agents_with_uncertainty_research`

Branch: `feature/orchestration-hypothesis-testing-experiment`

PR #1: **68+** **progress comments** documenting the full arc

This deck:

`experiments/orchestration_hypothesis_testing/presentation/supervisor_review.`